


Bayesian Methods for Machine Learning

Lecture 2 - Fundamentals of machine learning

Simon Leglaive

CentraleSupélec, 2020-2021



Today, we will focus on supervised learning.

Supervised learning

Consider an unknown joint probability distribution $p^*(\mathbf{x}, y)$.

Assume training data $\mathcal{D} = \{(\mathbf{x}_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}_{i=1, \dots, N}$ where

$$(\mathbf{x}_i, y_i) \sim p^*(\mathbf{x}, y).$$

- For instance \mathbf{x}_i is a p -dimensional input feature vector and y_i is a scalar label (e.g., a category or a real value).
- The training data is generated i.i.d.
- The training data can be of any finite size N .
- In general, we do not have any prior information about $p^*(\mathbf{x}, y)$.

Inference

Supervised learning is usually concerned with the two following inference problems:

- **Classification**: Given $\mathcal{D} = \{(\mathbf{x}_i, y_i) \in \mathcal{X} \times \mathcal{Y} = \mathbb{R}^p \times \{1, \dots, C\}\}_{i=1, \dots, N}$, we want to estimate for any **new** input \mathbf{x} :

$$\arg \max_y p(y|\mathbf{x}).$$


- **Regression**: Given $\mathcal{D} = \{(\mathbf{x}_i, y_i) \in \mathcal{X} \times \mathcal{Y} = \mathbb{R}^p \times \mathbb{R}\}_{i=1, \dots, N}$, we want to estimate for any **new** input \mathbf{x} :

$$\mathbb{E}_{p(y|\mathbf{x})} [y] = \int_{\mathcal{Y}} y p(y|\mathbf{x}) dy.$$


Or more generally, inference consists in computing the posterior distribution

$$p(y|\mathbf{x})$$

for any new \mathbf{x} (i.e. that was not part of the training dataset).



Classification consists in identifying
a decision boundary between objects of distinct classes.



Regression aims at estimating relationships among (usually continuous) variables.

Empirical risk minimization

Consider a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ produced by some learning algorithm. The predictions of this function can be evaluated through a loss

$$\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R},$$

such that $\ell(y, f(\mathbf{x})) \geq 0$ measures how close the prediction $f(\mathbf{x})$ from y is.

Examples of loss functions

Classification: $\ell(y, f(\mathbf{x})) = \mathbf{1}_{y \neq f(\mathbf{x})}$

Regression: $\ell(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$

Let \mathcal{F} denote the hypothesis space, i.e. the set of all functions f than can be produced by the chosen learning algorithm.

We are looking for a function $f \in \mathcal{F}$ with a small **expected risk** (or generalization error)

$$R(f) = \mathbb{E}_{p^*(\mathbf{x}, y)} [\ell(y, f(\mathbf{x}))] = \mathbb{E}_{p^*(\mathbf{x})} [\mathbb{E}_{p^*(y|\mathbf{x})} [\ell(y, f(\mathbf{x}))]] .$$

This means that for a given data generating distribution $p^*(\mathbf{x}, y)$ and for a given hypothesis space \mathcal{F} , the optimal model is

$$f^* = \arg \min_{f \in \mathcal{F}} R(f).$$

Unfortunately, since $p^*(\mathbf{x}, y)$ is unknown, the expected risk cannot be evaluated and the optimal model cannot be determined.

However, if we have i.i.d. training data $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1, \dots, N}$, we can compute an estimate, the **empirical risk** (or training error)

$$\hat{R}(f, \mathcal{D}) = \frac{1}{N} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}} \ell(y_i, f(\mathbf{x}_i)).$$

This estimate is **unbiased** and can be used for finding a good enough approximation of f^* . This results into the **empirical risk minimization principle**:


$$f_{\mathcal{D}}^* = \arg \min_{f \in \mathcal{F}} \hat{R}(f, \mathcal{D})$$

Most machine learning algorithms, including **neural networks**, implement empirical risk minimization.

Under regularity assumptions, empirical risk minimizers converge:

$$\lim_{N \rightarrow \infty} f_{\mathcal{D}}^* = f^*$$

Regression example





Consider the joint probability distribution $p^*(x, y)$ induced by the data generating process

$$(x, y) \sim p^*(x, y) \Leftrightarrow x \sim \mathcal{U}([-10; 10]), \epsilon \sim \mathcal{N}(0, \sigma^2), y = g(x) + \epsilon$$

where $x \in \mathbb{R}$, $y \in \mathbb{R}$ and g is an unknown polynomial of degree 3.

PREDICT TRAFFIC JAMS



REGRESSION

Regression is used to study the relationship between two continuous variables.

Of course, it can be extended to higher dimensions.

Step 1: Defining the model

Our goal is to find a function f that makes good predictions on average over $p^*(x, y)$.

Consider the hypothesis space $f \in \mathcal{F}$ of polynomials of degree 3 defined through their parameters $\mathbf{w} \in \mathbb{R}^4$ such that

$$\hat{y} \triangleq f(x; \mathbf{w}) = \sum_{d=0}^3 w_d x^d$$

Step 2: Defining the loss function

For this regression problem, we use the squared error loss

$$\ell(y, f(x; \mathbf{w})) = (y - f(x; \mathbf{w}))^2$$

to measure how wrong the predictions are.

Therefore, our goal is to find the best value \mathbf{w}^\star such

$$\begin{aligned}\mathbf{w}^\star &= \arg \min_{\mathbf{w}} R(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} \mathbb{E}_{p^\star(x,y)} [(y - f(x; \mathbf{w}))^2]\end{aligned}$$


Step 3: Training

Given a large enough training set $\mathcal{D} = \{(x_i, y_i) | i = 1, \dots, N\}$, the empirical risk minimization principle tells us that a good estimate $\mathbf{w}_{\mathcal{D}}^*$ of \mathbf{w}^* can be found by minimizing the empirical risk:

$$\begin{aligned}\mathbf{w}_{\mathcal{D}}^* &= \arg \min_{\mathbf{w}} \hat{R}(\mathbf{w}, \mathcal{D}) \\ &= \arg \min_{\mathbf{w}} \frac{1}{N} \sum_{(x_i, y_i) \in \mathcal{D}} (y_i - f(x_i; \mathbf{w}))^2 \\ &= \arg \min_{\mathbf{w}} \frac{1}{N} \sum_{(x_i, y_i) \in \mathcal{D}} \left(y_i - \sum_{d=0}^3 w_d x_i^d \right)^2 \\ &= \arg \min_{\mathbf{w}} \frac{1}{N} \left\| \underbrace{\begin{pmatrix} y_1 \\ y_2 \\ \dots \\ y_N \end{pmatrix}}_{\mathbf{y}} - \underbrace{\begin{pmatrix} x_1^0 & \dots & x_1^3 \\ x_2^0 & \dots & x_2^3 \\ \dots & \dots & \dots \\ x_N^0 & \dots & x_N^3 \end{pmatrix}}_{\mathbf{X}} \begin{pmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{pmatrix} \right\|^2\end{aligned}$$

This is **ordinary least squares** regression, for which the solution is known analytically:


$$\mathbf{w}_{\mathcal{D}}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$





In many situations, the problem is more difficult and we cannot find the solution analytically. We resort to **iterative optimization algorithms**, such as (variants of) gradient descent.


The expected risk minimizer $f(x; \mathbf{w}^*)$ within our hypothesis space \mathcal{F} (polynomials of degree 3) is $g(x)$ itself (i.e. the polynomial of degree 3 with the true parameters).


Therefore, on this toy problem, we can verify that $f(x; \mathbf{w}_{\mathcal{D}}^*) \rightarrow f(x; \mathbf{w}^*) = g(x)$ as $N \rightarrow \infty$.



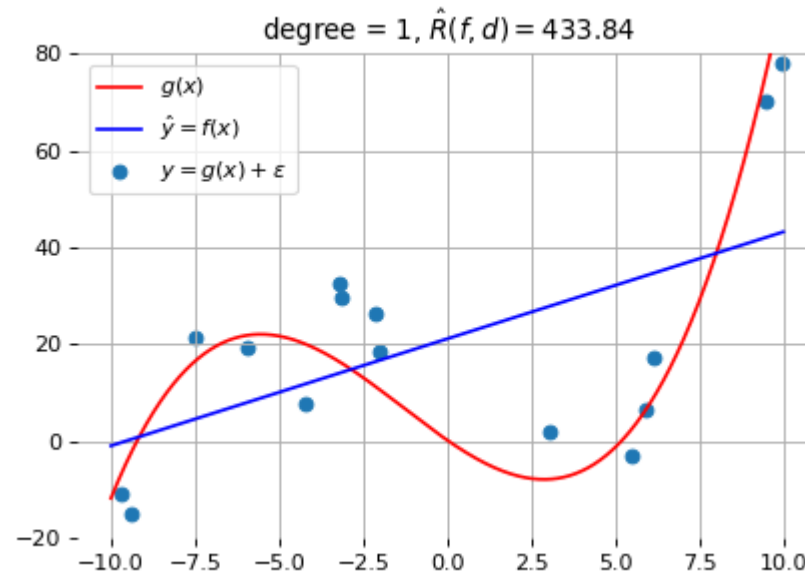






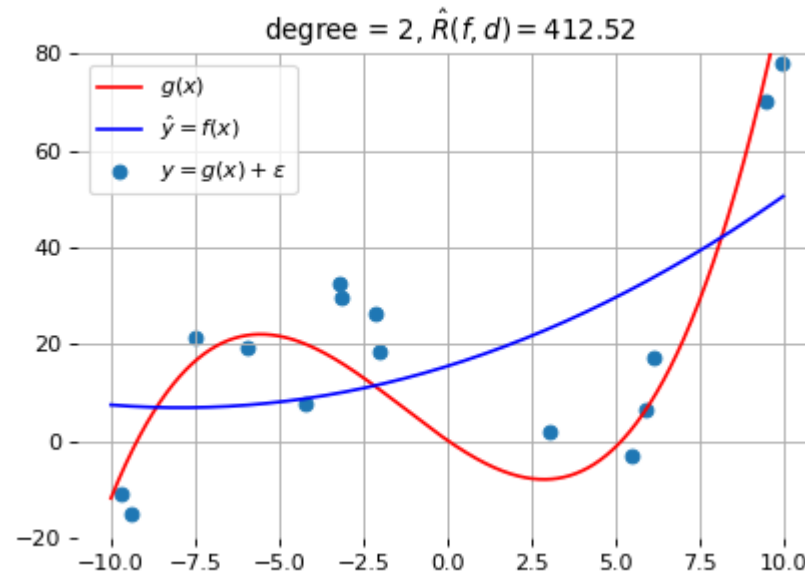


What if we consider a hypothesis space \mathcal{F} in which candidate functions f are either too "simple" or too "complex" with respect to the true data generating process?



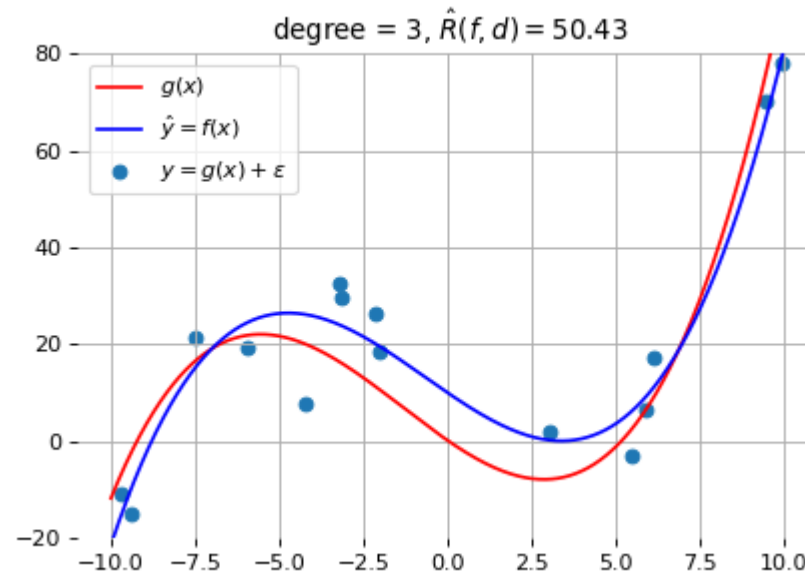
\mathcal{F} = polynomials of degree 1

What if we consider a hypothesis space \mathcal{F} in which candidate functions f are either too "simple" or too "complex" with respect to the true data generating process?



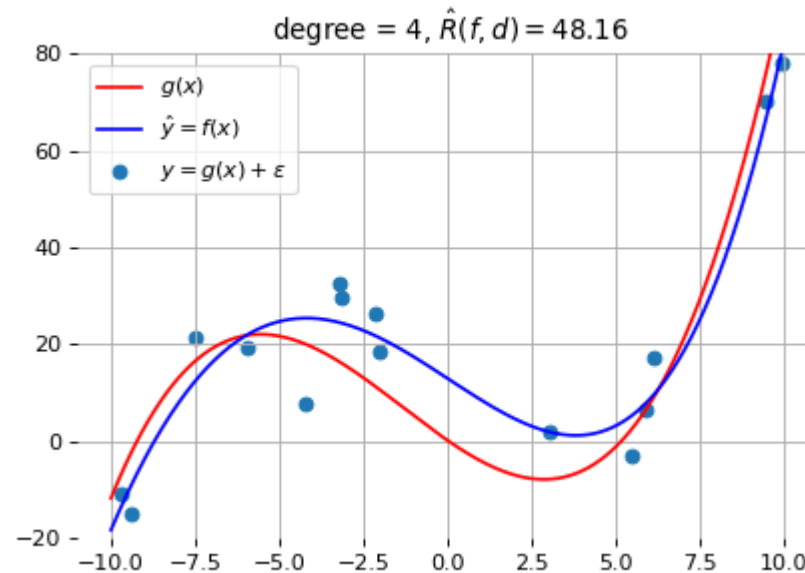
\mathcal{F} = polynomials of degree 2

What if we consider a hypothesis space \mathcal{F} in which candidate functions f are either too "simple" or too "complex" with respect to the true data generating process?



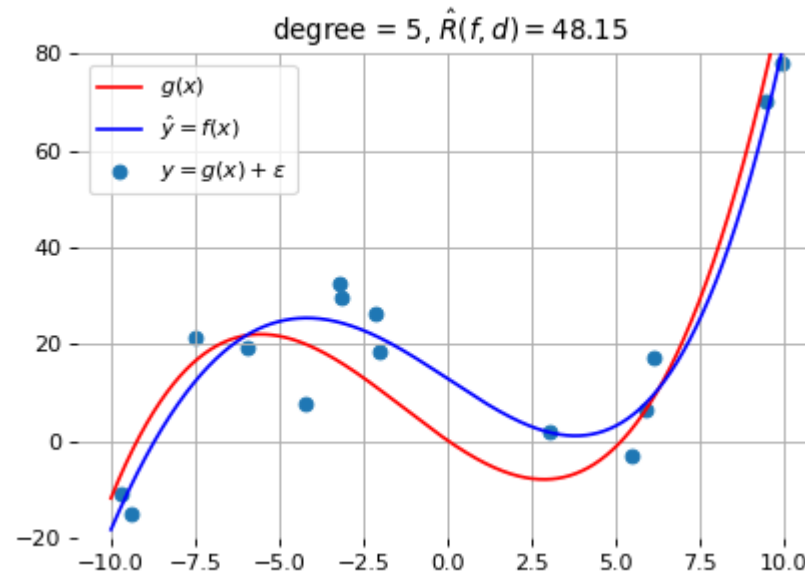
\mathcal{F} = polynomials of degree 3

What if we consider a hypothesis space \mathcal{F} in which candidate functions f are either too "simple" or too "complex" with respect to the true data generating process?



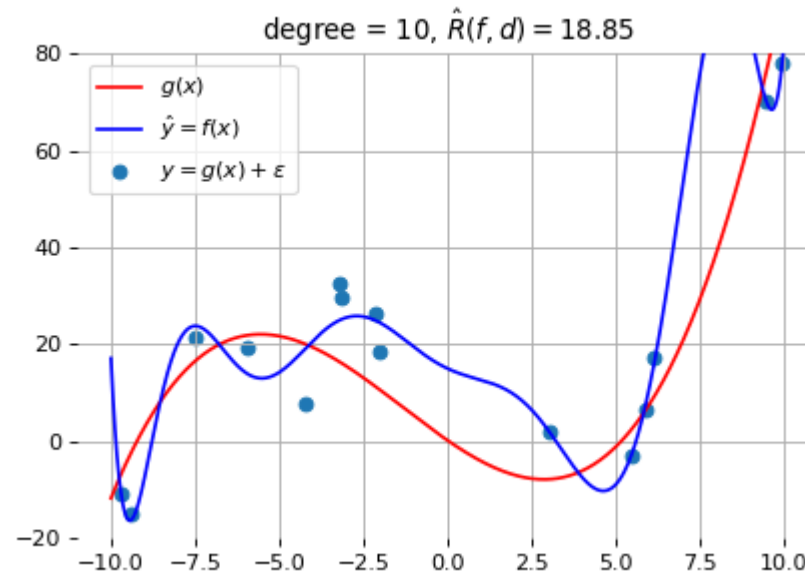
\mathcal{F} = polynomials of degree 4

What if we consider a hypothesis space \mathcal{F} in which candidate functions f are either too "simple" or too "complex" with respect to the true data generating process?




\mathcal{F} = polynomials of degree 5

What if we consider a hypothesis space \mathcal{F} in which candidate functions f are either too "simple" or too "complex" with respect to the true data generating process?



\mathcal{F} = polynomials of degree 10



Error vs. degree d of the polynomial.

Bayes risk and model

Let $\mathcal{Y}^{\mathcal{X}}$ be the set of all functions $f : \mathcal{X} \rightarrow \mathcal{Y}$.

We define the **Bayes risk** as the minimal expected risk over all possible functions,

$$R_B = \min_{f \in \mathcal{Y}^{\mathcal{X}}} R(f),$$

and call **Bayes model** the model f_B that achieves this minimum.

No model f can perform better than f_B .

The **capacity** of an hypothesis space \mathcal{F} induced by a learning algorithm intuitively represents the ability to find a good model $f \in \mathcal{F}$ that can fit any function, regardless of its complexity.

In practice, capacity can be controlled through hyper-parameters of the learning algorithm. For example:

- The degree of the family of polynomials;
- The number of layers in a neural network;
- The number of training iterations;
- Regularization terms.


Underfitting and overfitting

- If the capacity of \mathcal{F} is too low, then $f_B \notin \mathcal{F}$ and $R(f) - R_B$ is large for any $f \in \mathcal{F}$, including f^\star and $f_{\mathcal{D}}^\star$. Such models f are said to **underfit** the data.
- If the capacity of \mathcal{F} is too high, then $f_B \in \mathcal{F}$ or $R(f^\star) - R_B$ is small.
However, because of the high capacity of the hypothesis space, the empirical risk minimizer $f_{\mathcal{D}}^\star$ could fit the training data arbitrarily well such that

$$R(f_{\mathcal{D}}^\star) \geq R_B \geq \hat{R}(f_{\mathcal{D}}^\star, \mathcal{D}) \geq 0.$$

This indicates that the empirical risk $\hat{R}(f_{\mathcal{D}}^\star, \mathcal{D})$ is a poor estimator of the expected risk $R(f_{\mathcal{D}}^\star)$. In this situation, $f_{\mathcal{D}}^\star$ becomes too specialized with respect to the true data generating process, $f_{\mathcal{D}}^\star$ is said to **overfit** the data.

Therefore, our goal is to adjust the capacity of the hypothesis space such that the **expected risk** of the empirical risk minimizer (the generalization error) $R(f_{\mathcal{D}}^*)$ gets as low as possible, and not simply the **empirical risk** of the empirical risk minimizer (training error) $\hat{R}(f_{\mathcal{D}}^*, \mathcal{D})$.



An unbiased estimate of the expected risk can be obtained by evaluating $f_{\mathcal{D}}^*$ on data $\mathcal{D}_{\text{test}}$ independent from the training samples \mathcal{D} :


$$\hat{R}(f_{\mathcal{D}}^*, \mathcal{D}_{\text{test}}) = \frac{1}{N_{\text{test}}} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}_{\text{test}}} \ell(y_i, f_{\mathcal{D}}^*(\mathbf{x}_i))$$

This **test error** estimate can be used to evaluate the actual performance of the model. However, it should not be used, at the same time, for model selection.




Error vs. degree d of the polynomial.

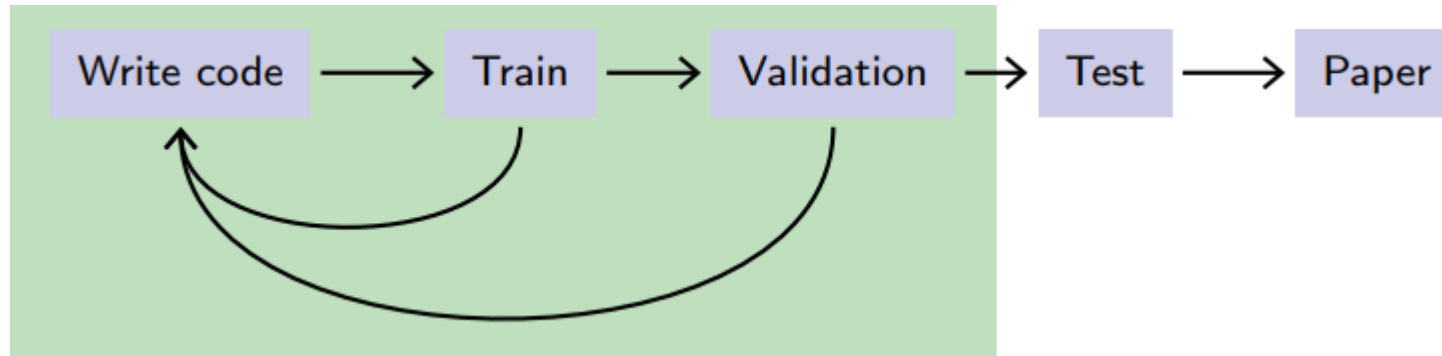
(Proper) evaluation protocol



There may be over-fitting, but it does not bias the final performance evaluation.



This should be **avoided** at all costs!



Instead, keep a separate validation set for tuning the hyper-parameters.

Bias-variance decomposition

Consider a **fixed point** x and the prediction $\hat{y} = f_{\mathcal{D}}^*(x)$ of the empirical risk minimizer at x .

Then the local expected risk of $f_{\mathcal{D}}^*$ is


$$\begin{aligned} R(f_{\mathcal{D}}^*|x) &= \mathbb{E}_{p^*(y|x)} [(y - f_{\mathcal{D}}^*(x))^2] \\ &= \mathbb{E}_{p^*(y|x)} [(y - f_B(x) + f_B(x) - f_{\mathcal{D}}^*(x))^2] \\ &= \mathbb{E}_{p^*(y|x)} [(y - f_B(x))^2] + \mathbb{E}_{p^*(y|x)} [(f_B(x) - f_{\mathcal{D}}^*(x))^2] \\ &= R(f_B|x) + (f_B(x) - f_{\mathcal{D}}^*(x))^2 \end{aligned}$$


where


- $R(f_B|x)$ is the local expected risk of the Bayes model. This term cannot be reduced.
- $(f_B(x) - f_{\mathcal{D}}^*(x))^2$ represents the discrepancy between f_B and $f_{\mathcal{D}}^*$.


Note that: $R(f) = \mathbb{E}_{p^*(\mathbf{x}, y)} [\ell(y, f(\mathbf{x}))] = \mathbb{E}_{p^*(\mathbf{x})} [\mathbb{E}_{p^*(y|\mathbf{x})} [\ell(y, f(\mathbf{x}))]] = \mathbb{E}_{p^*(\mathbf{x})} [R(f|x)]$


If \mathcal{D} is itself considered as a random variable, then $f_{\mathcal{D}}^*$ is also a random variable, along with its predictions \hat{y} .











Formally, the expected local expected risk yields to:

$$\begin{aligned} \mathbb{E}_{\mathcal{D}} [R(f_{\mathcal{D}}^*|x)] &= \mathbb{E}_{\mathcal{D}} [R(f_B|x) + (f_B(x) - f_{\mathcal{D}}^*(x))^2] \\ &= R(f_B|x) + \mathbb{E}_{\mathcal{D}} [(f_B(x) - f_{\mathcal{D}}^*(x))^2] \\ &= \underbrace{R(f_B|x)}_{\text{noise}(x)} + \underbrace{(f_B(x) - \mathbb{E}_{\mathcal{D}} [f_{\mathcal{D}}^*(x)])^2}_{\text{bias}^2(x)} + \underbrace{\mathbb{E}_{\mathcal{D}} [(\mathbb{E}_{\mathcal{D}} [f_{\mathcal{D}}^*(x)] - f_{\mathcal{D}}^*(x))^2]}_{\text{var}(x)} \end{aligned}$$

This decomposition is known as the **bias-variance** decomposition.

- The noise term quantity is the irreducible part of the expected risk.
- The bias term measures the discrepancy between the average model and the Bayes model.
- The variance term quantities the variability of the predictions.

Bias-variance trade-off

- Reducing the capacity makes $f_{\mathcal{D}}^*$ fit the data less on average, which increases the bias term.
- Increasing the capacity makes $f_{\mathcal{D}}^*$ vary a lot with the training data, which increases the variance term.

Maximum Likelihood and maximum a posteriori

Maximum Likelihood

Following the principle of empirical risk minimization, let $\mathcal{L}(\theta)$ denote a loss function defined over the model parameters by:

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}} \ell(y_i, f(\mathbf{x}_i; \theta)).$$

For both classification and regression, we can interpret $f(\mathbf{x}; \theta)$ as defining a model of the posterior distribution $p(y|\mathbf{x}; \theta)$.

Therefore, we can define the loss $\ell(\cdot, \cdot)$ as the negative log-likelihood (NLL):

$$\begin{aligned} \ell(y, f(\mathbf{x}; \theta)) &= -\ln p(\mathbf{x}, y; \theta) \\ &= -\ln p(y|\mathbf{x}; \theta) - \ln p(\mathbf{x}) \\ &= -\ln p(y|\mathbf{x}; \theta) + cst(\theta) \end{aligned}$$

Maximum a posteriori

We could also treat θ as a random variable and define the loss $\ell(\cdot, \cdot)$ as the negative log-posterior:

$$\begin{aligned}\ell(y, f(\mathbf{x}; \theta)) &= -\ln p(\theta | \mathbf{x}, y) \\ &= -\ln p(\mathbf{x}, y | \theta) - \ln p(\theta) + \text{cst}(\theta) \\ &= -\ln p(y | \mathbf{x}; \theta) - \ln p(\theta) + \text{cst}(\theta)\end{aligned}$$

It results in the negative log-likelihood plus a regularization term over θ which is the negative prior.

Binary classification

- **Training data:** $(\mathbf{x}, y) \in \mathcal{X} \times \mathcal{Y}$ with $\mathcal{X} = \mathbb{R}^p$ and $\mathcal{Y} = \{0, 1\}$.
- **Model:** $p(y = 1|\mathbf{x}; \theta) = f(\mathbf{x}; \theta)$ and $p(y = 0|\mathbf{x}; \theta) = 1 - f(\mathbf{x}; \theta)$.

It can be compactly rewritten as follows for all $y \in \mathcal{Y}$:

$$p(y|\mathbf{x}; \theta) = \left(f(\mathbf{x}; \theta)\right)^y \left(1 - f(\mathbf{x}; \theta)\right)^{(1-y)}.$$

- **Constraint:** $f(\mathbf{x}; \theta) \in [0, 1]$.
- The **NLL** gives the **binary cross-entropy** loss:

$$\begin{aligned}\ell(y, f(\mathbf{x}; \theta)) &= -\ln p(\mathbf{x}, y; \theta) \\ &= -\ln p(y|\mathbf{x}; \theta) - \ln p(\mathbf{x}) \\ &= -y \ln \left(f(\mathbf{x}; \theta)\right) - (1 - y) \ln \left(1 - f(\mathbf{x}; \theta)\right) + cst(\theta)\end{aligned}$$

C-class classification

- **Training data:** $(\mathbf{x}, y) \in \mathcal{X} \times \mathcal{Y}$ with $\mathcal{X} = \mathbb{R}^p$ and $\mathcal{Y} = \{1, \dots, C\}$.
- **Model:** $p(y = c | \mathbf{x}; \theta) = f_c(\mathbf{x}; \theta)$ for all $c \in \{1, \dots, C\}$.

It can be compactly rewritten as follows for all $y \in \mathcal{Y}$:

$$p(y | \mathbf{x}; \theta) = \prod_{c=1}^C p(y = c | \mathbf{x}; \theta)^{\mathbf{1}_{y=c}} = \prod_{c=1}^C f_c(\mathbf{x}; \theta)^{\mathbf{1}_{y=c}}.$$

- **Constraint:** $f(\mathbf{x}; \theta) \in [0, 1]^C$ and $\sum_{c=1}^C f_c(\mathbf{x}; \theta) = 1$ where $f_c(\mathbf{x}; \theta)$ is the c -th entry of $f(\mathbf{x}; \theta)$.
- The **NLL** gives the **cross-entropy** loss:

$$\ell(y, f(\mathbf{x}; \theta)) = - \sum_{c=1}^C \mathbf{1}_{y=c} \ln \left(f_c(\mathbf{x}; \theta) \right) + cst(\theta).$$

Regression

- **Training data:** $(\mathbf{x}, \mathbf{y}) \in \mathcal{X} \times \mathcal{Y}$ with $\mathcal{X} = \mathbb{R}^p$ and $\mathcal{Y} = \mathbb{R}^q$.
- **Model:** $p(\mathbf{y}|\mathbf{x}; \theta) = \mathcal{N}(\mathbf{y}; f(\mathbf{x}; \theta), \mathbf{I}) = (2\pi)^{-q/2} \exp \left(-\frac{1}{2} \|\mathbf{y} - f(\mathbf{x}; \theta)\|_2^2 \right)$.
- **Constraint:** $f(\mathbf{x}; \theta) \in \mathbb{R}^q$.
- The **NLL** gives the **squared error** loss:

$$\ell(y, f(\mathbf{x}; \theta)) = \frac{1}{2} \|\mathbf{y} - f(\mathbf{x}; \theta)\|_2^2 + cst(\theta).$$

Lab session on multinomial logistic regression

